Azure Banking Data Platform

A production-grade, real-time data engineering pipeline built using **Azure Event-Driven Architecture**, **Cosmos DB**, and the **Databricks Lakehouse Platform**.

---

1. Project Overview

This solution delivers an end-to-end banking data platform capable of handling **real-time ATM and UPI transactions**, **customer profile updates**, **KYC records**, and **branch performance analytics**.

The platform integrates:

- **Event-driven ingestion** using Event Grid, Azure Functions, and Service Bus
- **Operational data store** using Azure Cosmos DB

- **Data Lakehouse ETL** using Databricks (Bronze → Silver → Gold)
- **SCD-2 dimension management**
- **SQL Data Warehouse serving layer**
- **Power BI dashboards for Customer 360, Fraud Metrics, Branch Performance**
- **CI/CD workflow using GitHub Actions / Azure DevOps**

---

2. High-Level Architecture

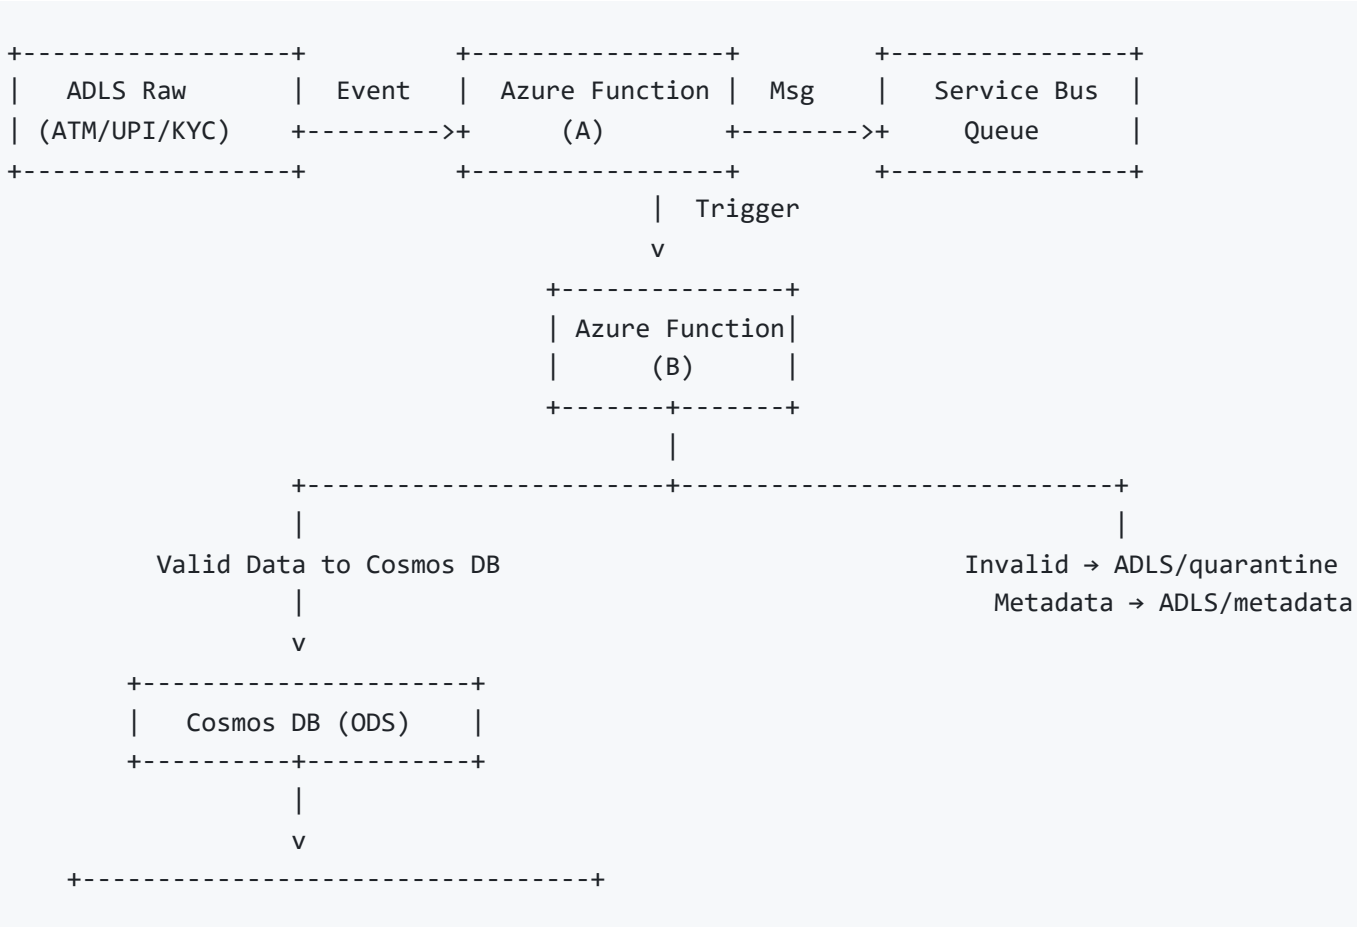

```
|     Databricks Lakehouse        |
|  Bronze → Silver → Gold Layers  |
+-----------------+---------------+
                  |
                  v
          +-----------------+
          | SQL Data Ware-  |
          | house (Serving) |
          +-------+---------+
                  |
                  v
          +----------------+
          |   Power BI     |
          | Dashboards     |
          +----------------+
```

3. Azure Resources Used

Service	Purpose
ADLS Gen2	Raw, Quarantine, Metadata, Bronze, Silver, Gold zones
Event Grid	Detects new file uploads in ADLS
Azure Function A	Receives EventGrid events and pushes file metadata into Service Bus
Service Bus Queue	Reliable buffer for downstream processing
Azure Function B (Queue Trigger)	Validates payload → writes valid records to Cosmos DB / invalid to ADLS
Cosmos DB	Operational store for ATM, UPI, and Customer data

Service	Purpose
Databricks	Orchestrates Bronze → Silver → Gold ETL
Azure SQL DB / SQL DWH	Serving layer for BI tools
Power BI	Analytics dashboards
CI/CD (GitHub Actions)	Notebook sync + Function App deployment

---

4. ADLS Folder Structure

```
/raw
    /atm
    /upi
    /customers
    /kyc

/quarantine
/metadata

/etl
    /bronze
    /silver
    /gold
```

---

5. Ingestion Workflow

Step 1 — File lands in ADLS (raw)

ATM, UPI, Customer, KYC files are uploaded into `/raw`.

Step 2 — Event Grid triggers Function A

Detects `BlobCreated` events and extracts:

- file_url
- file_name
- source_type (ATM/UPI/CUSTOMER/KYC)

Step 3 — Function A → Event Grid Trigger

Minimal enrichment and routing.

Step 4 — Function B (Queue-trigger)

Validates schema, datatype, required fields.

- **If valid:** → Insert into Cosmos DB → Store metadata in `/metadata`
- **If invalid:** → Store full payload in `/quarantine`

Step 5 — Databricks Ingestion

- Reads **Cosmos DB** for transactions & customer profiles
- Reads **KYC directly from ADLS**

---

6. Databricks ETL (Lakehouse)

6.1 Bronze Layer

- Raw ingestion from Cosmos DB
- KYC raw ingestion from ADLS
- Light parsing
- No transformations
- Stored as Delta tables

6.2 Silver Layer

- Standardized schema
- Data quality cleanup
- Normalized structures
- Reference cleanup
- Joins across sources
- Error-handled and type-safe

6.3 Gold Layer (Final Analytics Layer)

Dimensions

- `dim_customer`
- `dim_account`
- `dim_kyc`

- `dim_date`

SCD-2 Dimensions

- `dim_customer_scd2`
- `dim_account_scd2`

Facts

- `fact_transactions`
- `fact_customer_profile`

Aggregates

- Daily transaction summary
- Customer monthly spending
- Branch performance
- Channel-wise distribution

---

7. Serving Layer (SQL Data Warehouse)

Gold tables and aggregates are written to **Azure SQL DB / SQL Warehouse**, enabling optimized BI reporting.

---

8. Power BI Dashboards

Dashboards supported:

Example placeholders:

9. CI/CD

CI Pipeline

- Linting & quality checks
- Validate Python Function code
- Validate Databricks notebooks
- Package Function App

CD Pipeline

- Deploy Azure Function App
- Sync Databricks notebooks (via Databricks CLI)
- Optionally trigger Databricks job
- Validate Cosmos DB + ADLS connectivity

Pipeline can be implemented in:

- GitHub Actions
- Azure DevOps Pipelines

---

10. Repository Structure

```
/src
    /function_app
    /notebooks
         /bronze
         /silver
         /gold
    /powerbi
```

```
    /sql
/docs
    /images
    architecture.md
/readme.md
```

---

11. How to Run the Project

Prerequisites

- Azure subscription
- Cosmos DB + ADLS + Function App + Service Bus
- Databricks Workspace
- SQL Database

Execution Steps

1. Upload sample data into `/raw` in ADLS
2. Event Grid → Functions → Service Bus → Cosmos DB automatically ingests
3. Databricks Jobs run Bronze → Silver → Gold pipeline
4. Gold tables are published to SQL DWH
5. Power BI connects to SQL DWH for reporting
6. CI/CD handles notebook + function deployments

---

12. Screenshot Placeholders

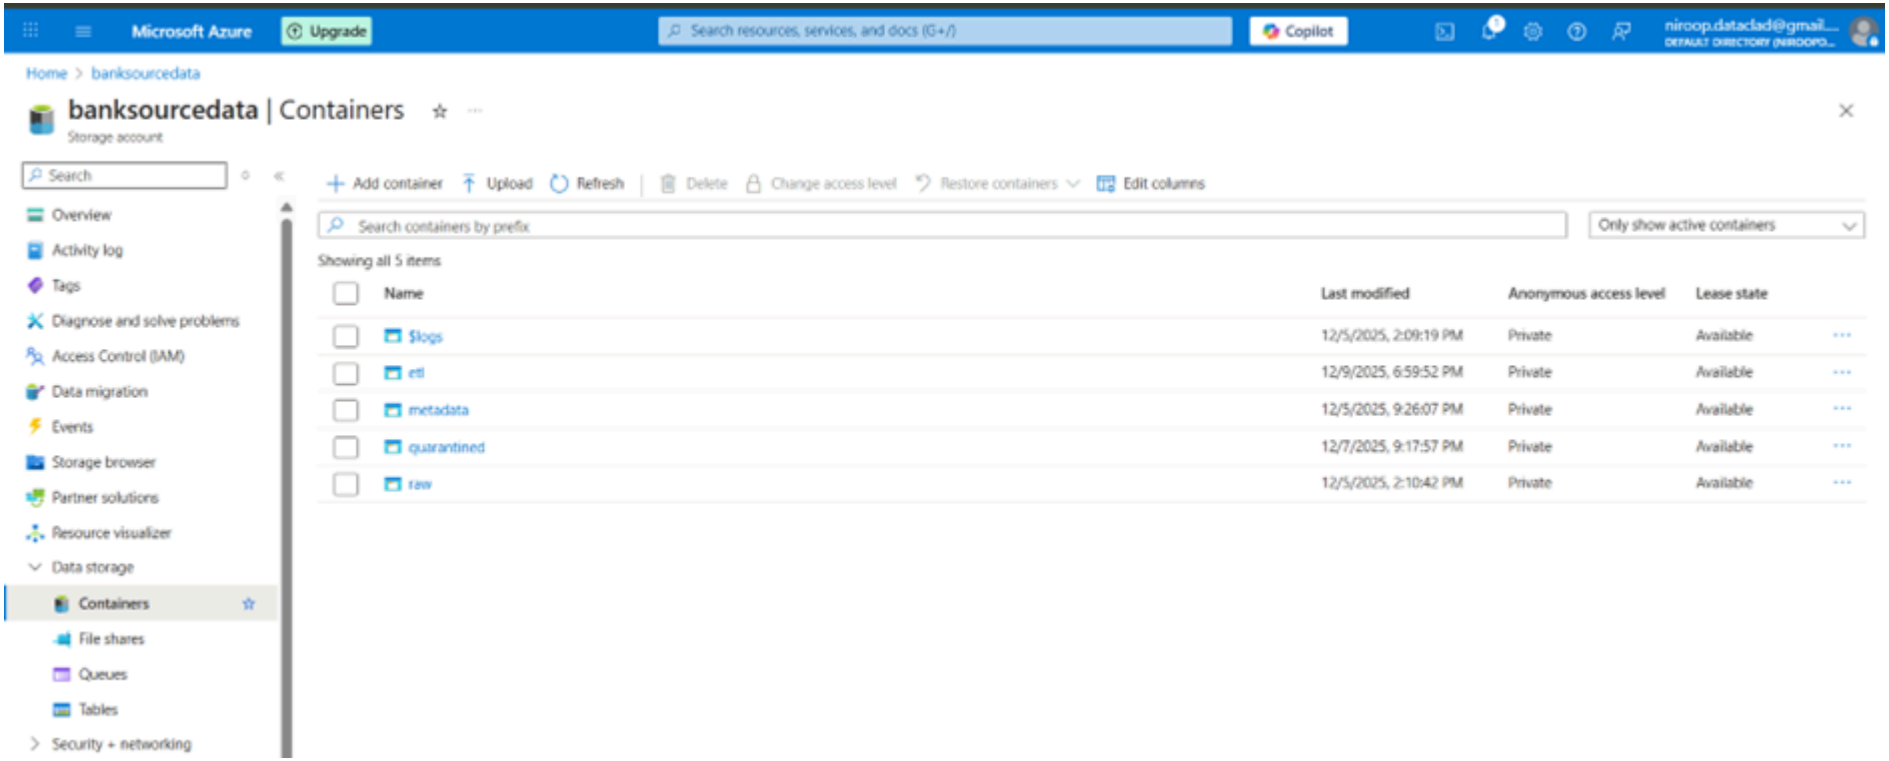
Microsoft Azure
Upgrade
Search resources, services, and docs (G+/)
Copilot
Home > banksourcedata
banksourcedata | Containers
Storage account
Search
Overview
Activity log
Tags
Diagnose and solve problems
Access Control (IAM)
Data migration
Events
Storage browser
Partner solutions
Resource visualizer
Data storage
Containers
File shares
Queues
Tables
Security + networking
Add container
Upload
Refresh
Delete
Change access level
Restore containers
Edit columns
Search containers by prefix
Only show active containers
Showing all 5 items
Name
Last modified
Anonymous access level
Lease state
$logs
12/5/2025, 2:09:19 PM
Private
Available
etl
12/9/2025, 6:59:52 PM
Private
Available
metadata
12/5/2025, 9:26:07 PM
Private
Available
quarantined
12/7/2025, 9:17:57 PM
Private
Available
raw
12/5/2025, 2:10:42 PM
Private
Available

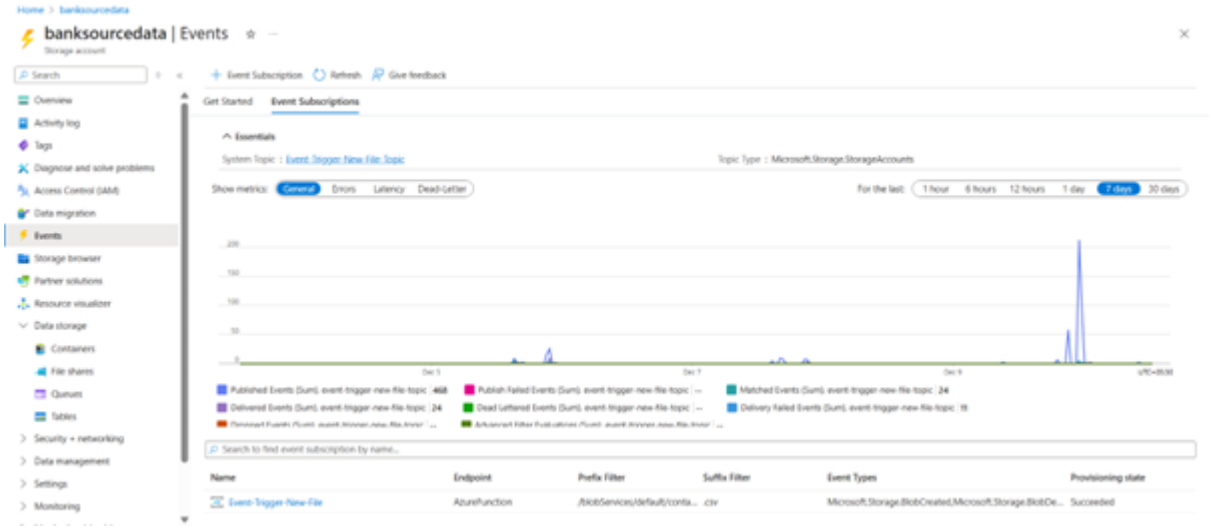
banksourcedata | Events
Storage account
Event Subscription
Refresh
Give feedback
Get Started
Event Subscriptions
Essentials
Show metrics:
General
Errors
Latency
Dead-Letter
For the last:
1 hour
6 hours
12 hours
1 day
7 days
30 days
Search to find event subscription by name...
Name
Endpoint
Prefix Filter
Suffix Filter
Event Types
Provisioning state
Event-Trigger-New-File
AzureFunction
.csv
Succeeded

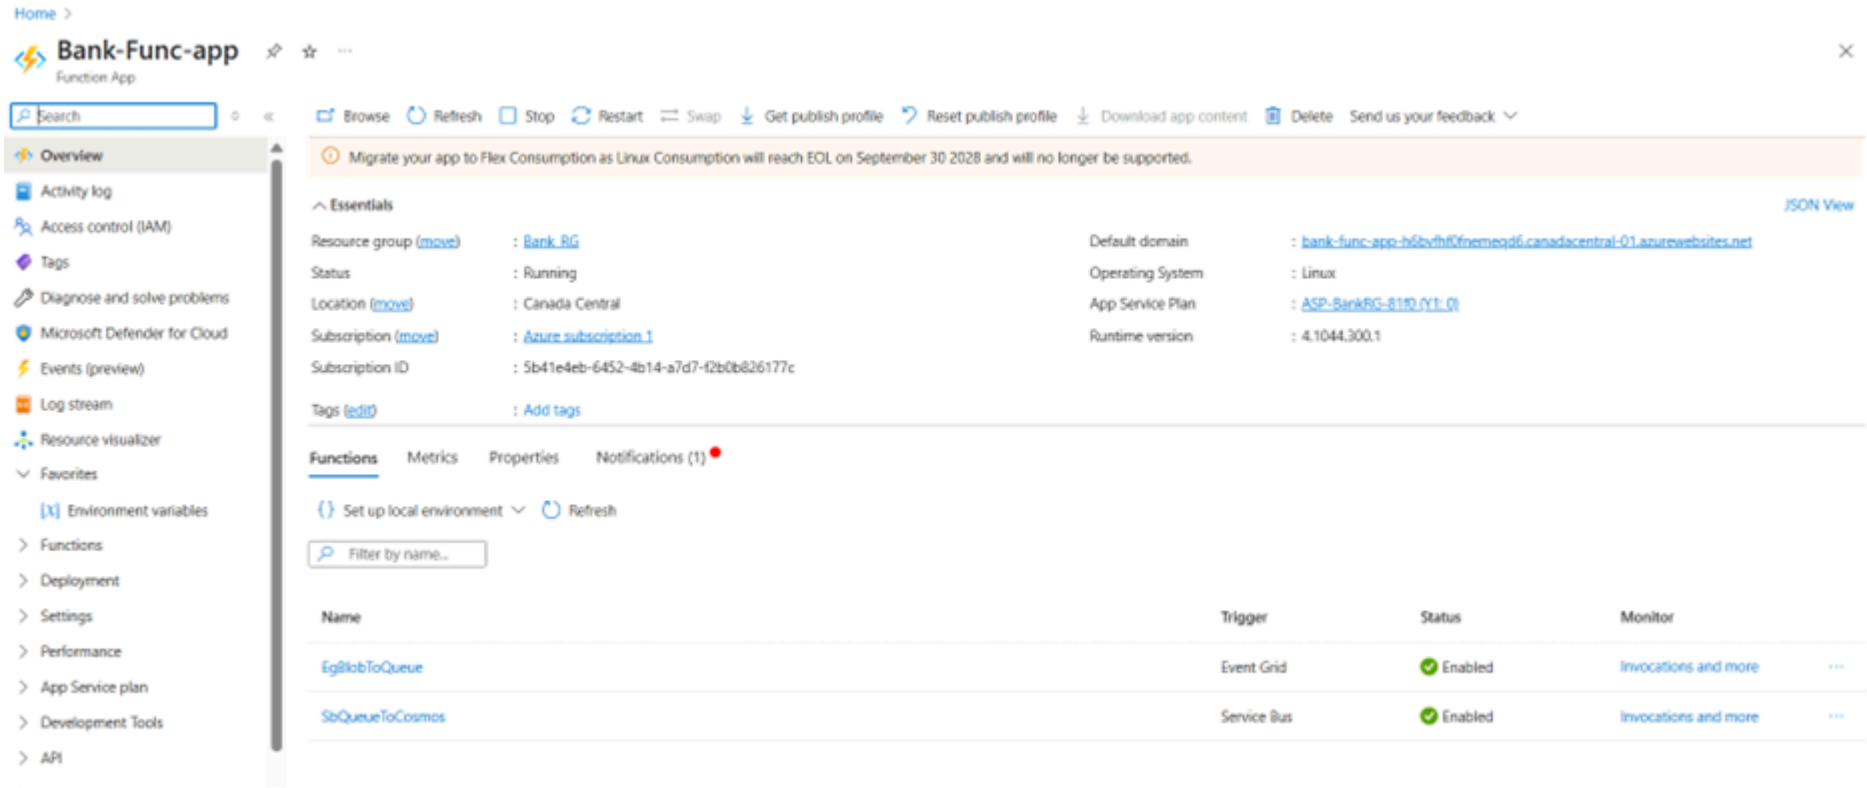
Home
Bank-Func-app
Function App
Search
Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems
Microsoft Defender for Cloud
Events (preview)
Log stream
Resource visualizer
Favorites
Environment variables
Functions
Deployment
Settings
Performance
App Service plan
Development Tools
API
Browse
Refresh
Stop
Restart
Swap
Get publish profile
Reset publish profile
Download app content
Delete
Send us your feedback
Migrate your app to Flex Consumption as Linux Consumption will reach EOL on September 30 2028 and will no longer be supported.
Essentials
JSON View
Resource group (move) : Bank_RG
Status : Running
Location (move) : Canada Central
Subscription (move) : Azure subscription 1
Subscription ID : 5b41e4eb-6452-4b14-a7d7-f2b0b826177c
Tags (edit) : Add tags
Default domain : bank-func-app-h6bvfhf0fnemeqd6.canadacentral-01.azurewebsites.net
Operating System : Linux
App Service Plan : ASP-BankRG-81f0 (Y1: 0)
Runtime version : 4.1044.300.1
Functions
Metrics
Properties
Notifications (1)
Set up local environment
Refresh
Filter by name...
Name
Trigger
Status
Monitor
EgBlobToQueue
Event Grid
Enabled
Invocations and more
SbQueueToCosmos
Service Bus
Enabled
Invocations and more

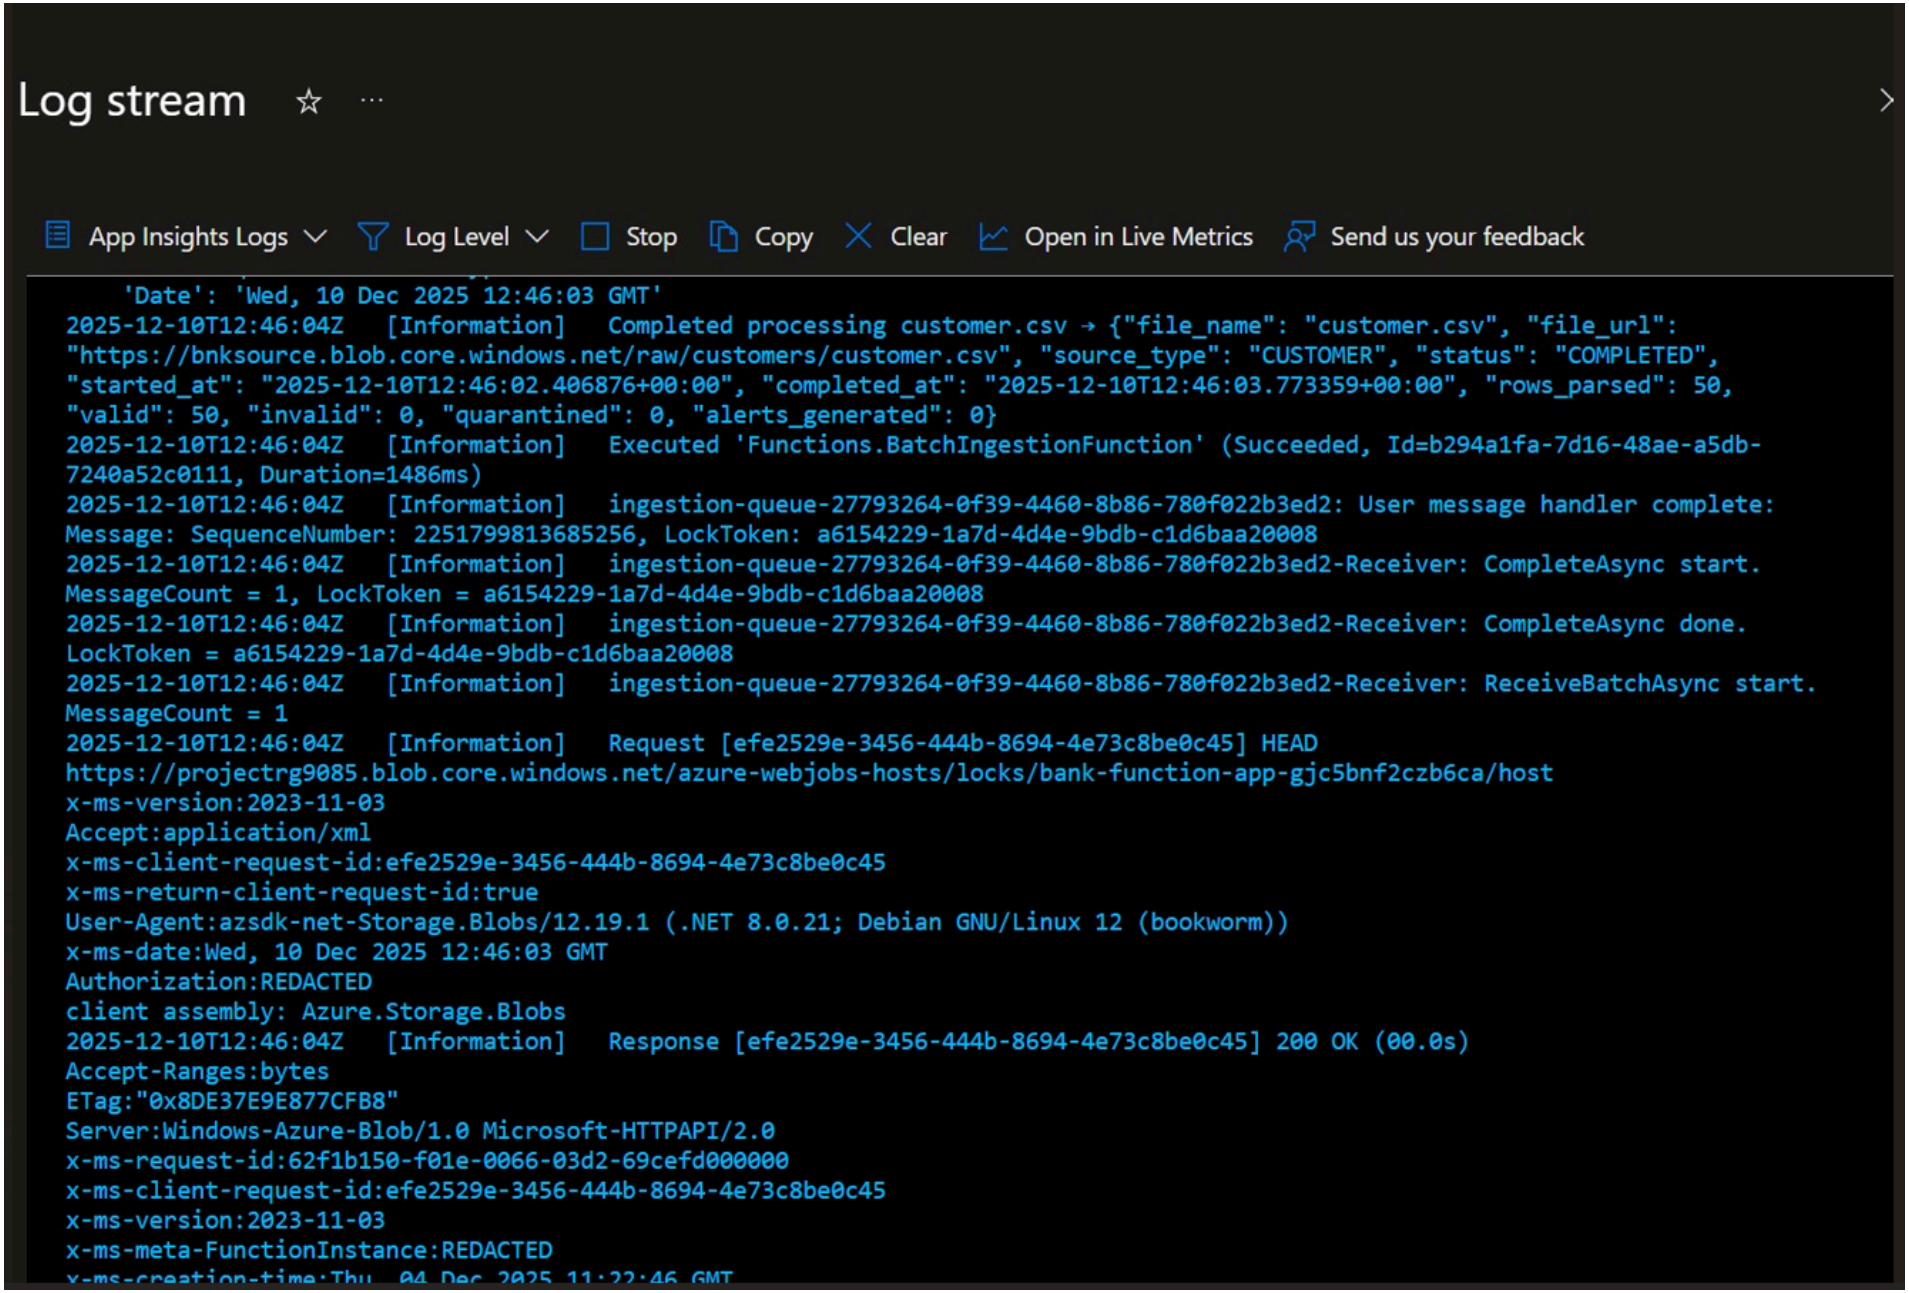
Log stream
App Insights Logs
Log Level
Stop
Copy
Clear
Open in Live Metrics
Send us your feedback
'Date': 'Wed, 10 Dec 2025 12:46:03 GMT'
2025-12-10T12:46:04Z [Information] Completed processing customer.csv → {"file_name": "customer.csv", "file_url": "https://bnksource.blob.core.windows.net/raw/customers/customer.csv", "source_type": "CUSTOMER", "status": "COMPLETED", "started_at": "2025-12-10T12:46:02.406876+00:00", "completed_at": "2025-12-10T12:46:03.773359+00:00", "rows_parsed": 50, "valid": 50, "invalid": 0, "quarantined": 0, "alerts_generated": 0}
2025-12-10T12:46:04Z [Information] Executed 'Functions.BatchIngestionFunction' (Succeeded, Id=b294a1fa-7d16-48ae-a5db-7240a52c0111, Duration=1486ms)
2025-12-10T12:46:04Z [Information] ingestion-queue-27793264-0f39-4460-8b86-780f022b3ed2: User message handler complete: Message: SequenceNumber: 2251799813685256, LockToken: a6154229-1a7d-4d4e-9bdb-c1d6baa20008
2025-12-10T12:46:04Z [Information] ingestion-queue-27793264-0f39-4460-8b86-780f022b3ed2-Receiver: CompleteAsync start. MessageCount = 1, LockToken = a6154229-1a7d-4d4e-9bdb-c1d6baa20008
2025-12-10T12:46:04Z [Information] ingestion-queue-27793264-0f39-4460-8b86-780f022b3ed2-Receiver: CompleteAsync done. LockToken = a6154229-1a7d-4d4e-9bdb-c1d6baa20008
2025-12-10T12:46:04Z [Information] ingestion-queue-27793264-0f39-4460-8b86-780f022b3ed2-Receiver: ReceiveBatchAsync start. MessageCount = 1
2025-12-10T12:46:04Z [Information] Request [efe2529e-3456-444b-8694-4e73c8be0c45] HEAD https://projectrg9085.blob.core.windows.net/azure-webjobs-hosts/locks/bank-function-app-gjc5bnf2czb6ca/host
x-ms-version:2023-11-03
Accept:application/xml
x-ms-client-request-id:efe2529e-3456-444b-8694-4e73c8be0c45
x-ms-return-client-request-id:true
User-Agent:azsdk-net-Storage.Blobs/12.19.1 (.NET 8.0.21; Debian GNU/Linux 12 (bookworm))
x-ms-date:Wed, 10 Dec 2025 12:46:03 GMT
Authorization:REDACTED
client assembly: Azure.Storage.Blobs
2025-12-10T12:46:04Z [Information] Response [efe2529e-3456-444b-8694-4e73c8be0c45] 200 OK (00.0s)
Accept-Ranges:bytes
ETag:"0x8DE37E9E877CFB8"
Server:Windows-Azure-Blob/1.0 Microsoft-HTTPAPI/2.0
x-ms-request-id:62f1b150-f01e-0066-03d2-69cefd000000
x-ms-client-request-id:efe2529e-3456-444b-8694-4e73c8be0c45
x-ms-version:2023-11-03
x-ms-meta-FunctionInstance:REDACTED

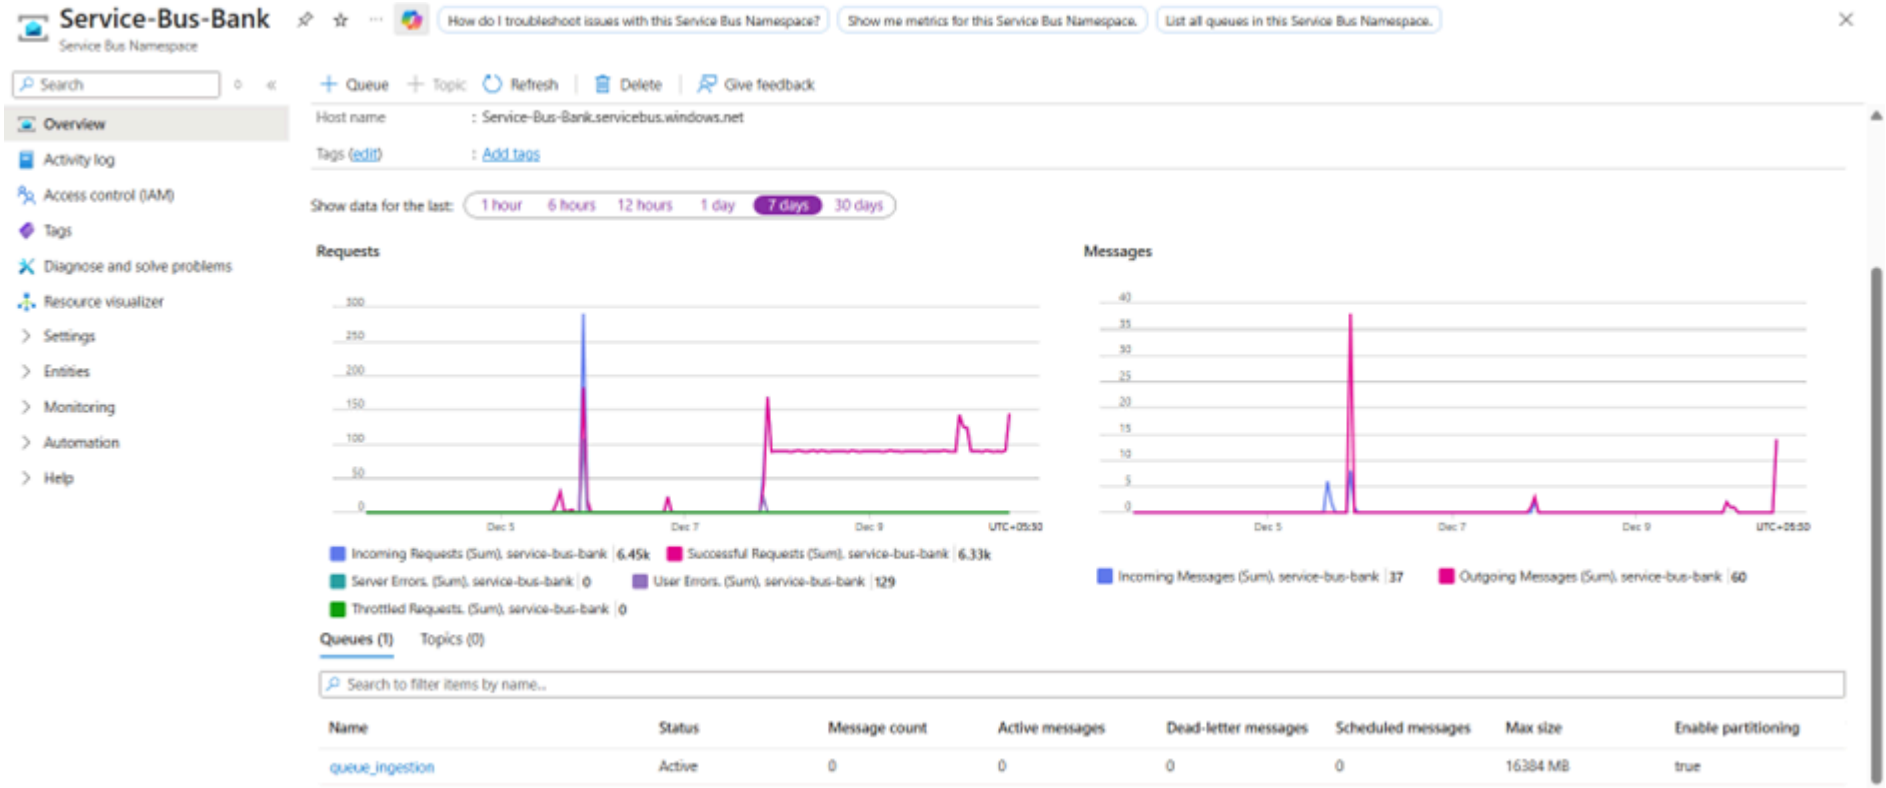
Service-Bus-Bank
Service Bus Namespace
How do I troubleshoot issues with this Service Bus Namespace?
Show me metrics for this Service Bus Namespace.
List all queues in this Service Bus Namespace.
Search
Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems
Resource visualizer
Settings
Entities
Monitoring
Automation
Help
Queue
Topic
Refresh
Delete
Give feedback
Host name : Service-Bus-Bank.servicebus.windows.net
Tags (edit) : Add tags
Show data for the last: 1 hour 6 hours 12 hours 1 day 7 days 30 days
Requests
Messages
Dec 5
Dec 7
Dec 9
UTC+05:30
Incoming Requests (Sum), service-bus-bank 6.45k
Successful Requests (Sum), service-bus-bank 6.33k
Server Errors (Sum), service-bus-bank 0
User Errors (Sum), service-bus-bank 129
Throttled Requests (Sum), service-bus-bank 0
Incoming Messages (Sum), service-bus-bank 37
Outgoing Messages (Sum), service-bus-bank 60
Queues (1)
Topics (0)
Search to filter items by name...
Name
Status
Message count
Active messages
Dead-letter messages
Scheduled messages
Max size
Enable partitioning
queue_ingestion
Active
0
0
0
0
16384 MB
true

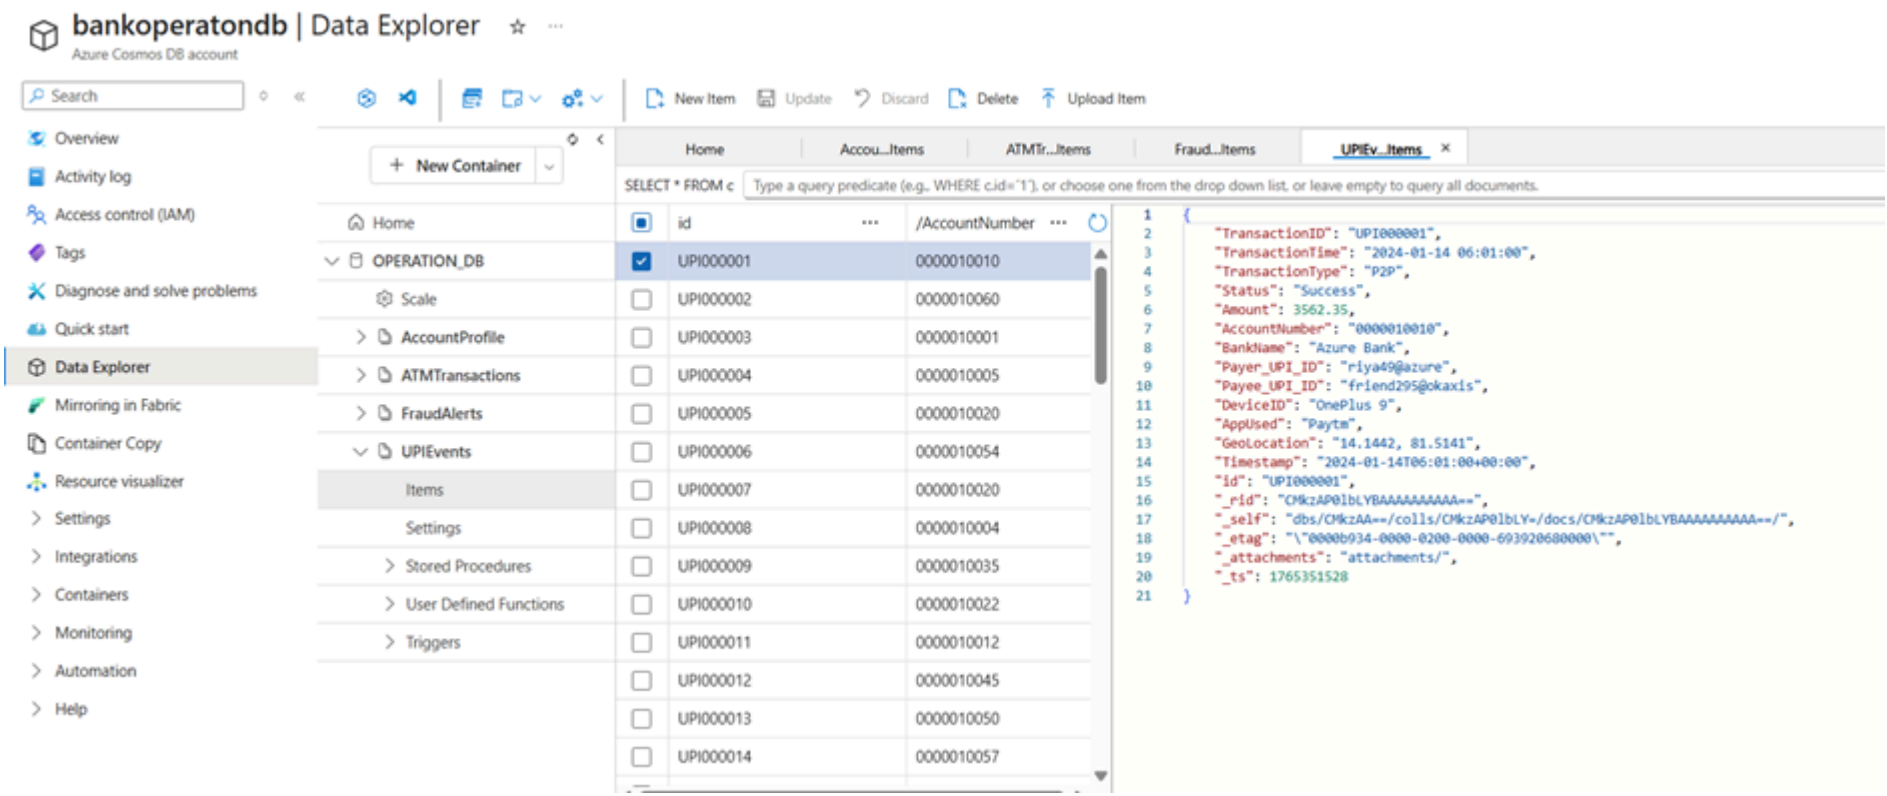
bankoperatondb | Data Explorer
Azure Cosmos DB account
Search
Overview
Activity log
Access control (IAM)
Tags
Diagnose and solve problems
Quick start
Data Explorer
Mirroring in Fabric
Container Copy
Resource visualizer
Settings
Integrations
Containers
Monitoring
Automation
Help
New Item
Update
Discard
Delete
Upload Item
New Container
Home
OPERATION_DB
Scale
AccountProfile
ATMTransactions
FraudAlerts
UPIEvents
Items
Settings
Stored Procedures
User Defined Functions
Triggers
Accou...Items
ATMTr...Items
Fraud...Items
UPIEv...Items
SELECT * FROM c
Type a query predicate (e.g., WHERE c.id='1'), or choose one from the drop down list, or leave empty to query all documents.
id
/AccountNumber
UPI000001 0000010010
UPI000002 0000010060
UPI000003 0000010001
UPI000004 0000010005
UPI000005 0000010020
UPI000006 0000010054
UPI000007 0000010020
UPI000008 0000010004
UPI000009 0000010035
UPI000010 0000010022
UPI000011 0000010012
UPI000012 0000010045
UPI000013 0000010050
UPI000014 0000010057
{
"TransactionID": "UPI000001",
"TransactionTime": "2024-01-14 06:01:00",
"TransactionType": "P2P",
"Status": "Success",
"Amount": 3562.35,
"AccountNumber": "0000010010",
"BankName": "Azure Bank",
"Payer_UPI_ID": "riya49@azure",
"Payee_UPI_ID": "friend295@okaxis",
"DeviceID": "OnePlus 9",
"AppUsed": "Paytm",
"GeoLocation": "14.1442, 81.5141",
"Timestamp": "2024-01-14T06:01:00+00:00",
"id": "UPI000001",
"_rid": "CMkzAP0lbLYBAAAAAAAAAA==",
"_self": "dbs/CMkzAA==/colls/CMkzAP0lbLY=/docs/CMkzAP0lbLYBAAAAAAAAAA==/",
"_etag": "\"0000b934-0000-0200-0000-693920680000\"",
"_attachments": "attachments/",
"_ts": 1765351528
}

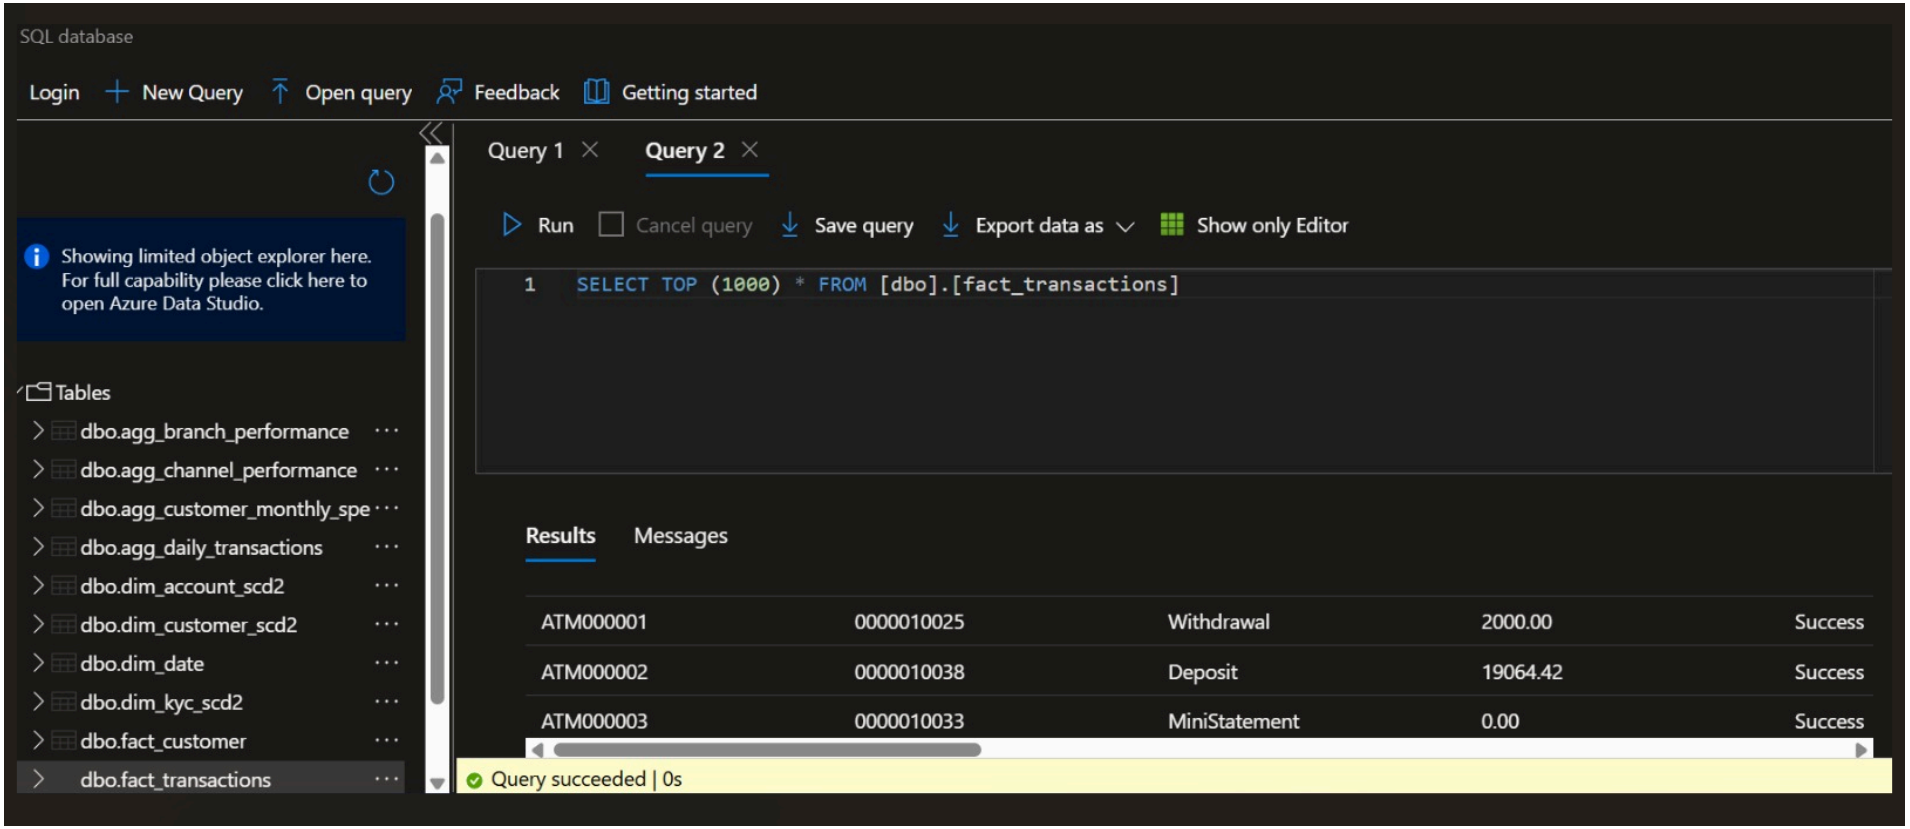
SQL database
Login
New Query
Open query
Feedback
Getting started
Query 1
Query 2
Run
Cancel query
Save query
Export data as
Show only Editor
Showing limited object explorer here. For full capability please click here to open Azure Data Studio.
Tables
dbo.agg_branch_performance
dbo.agg_channel_performance
dbo.agg_customer_monthly_spe
dbo.agg_daily_transactions
dbo.dim_account_scd2
dbo.dim_customer_scd2
dbo.dim_date
dbo.dim_kyc_scd2
dbo.fact_customer
dbo.fact_transactions
1 SELECT TOP (1000) * FROM [dbo].[fact_transactions]
Results
Messages
ATM000001 0000010025 Withdrawal 2000.00 Success
ATM000002 0000010038 Deposit 19064.42 Success
ATM000003 0000010033 MiniStatement 0.00 Success
Query succeeded | 0s

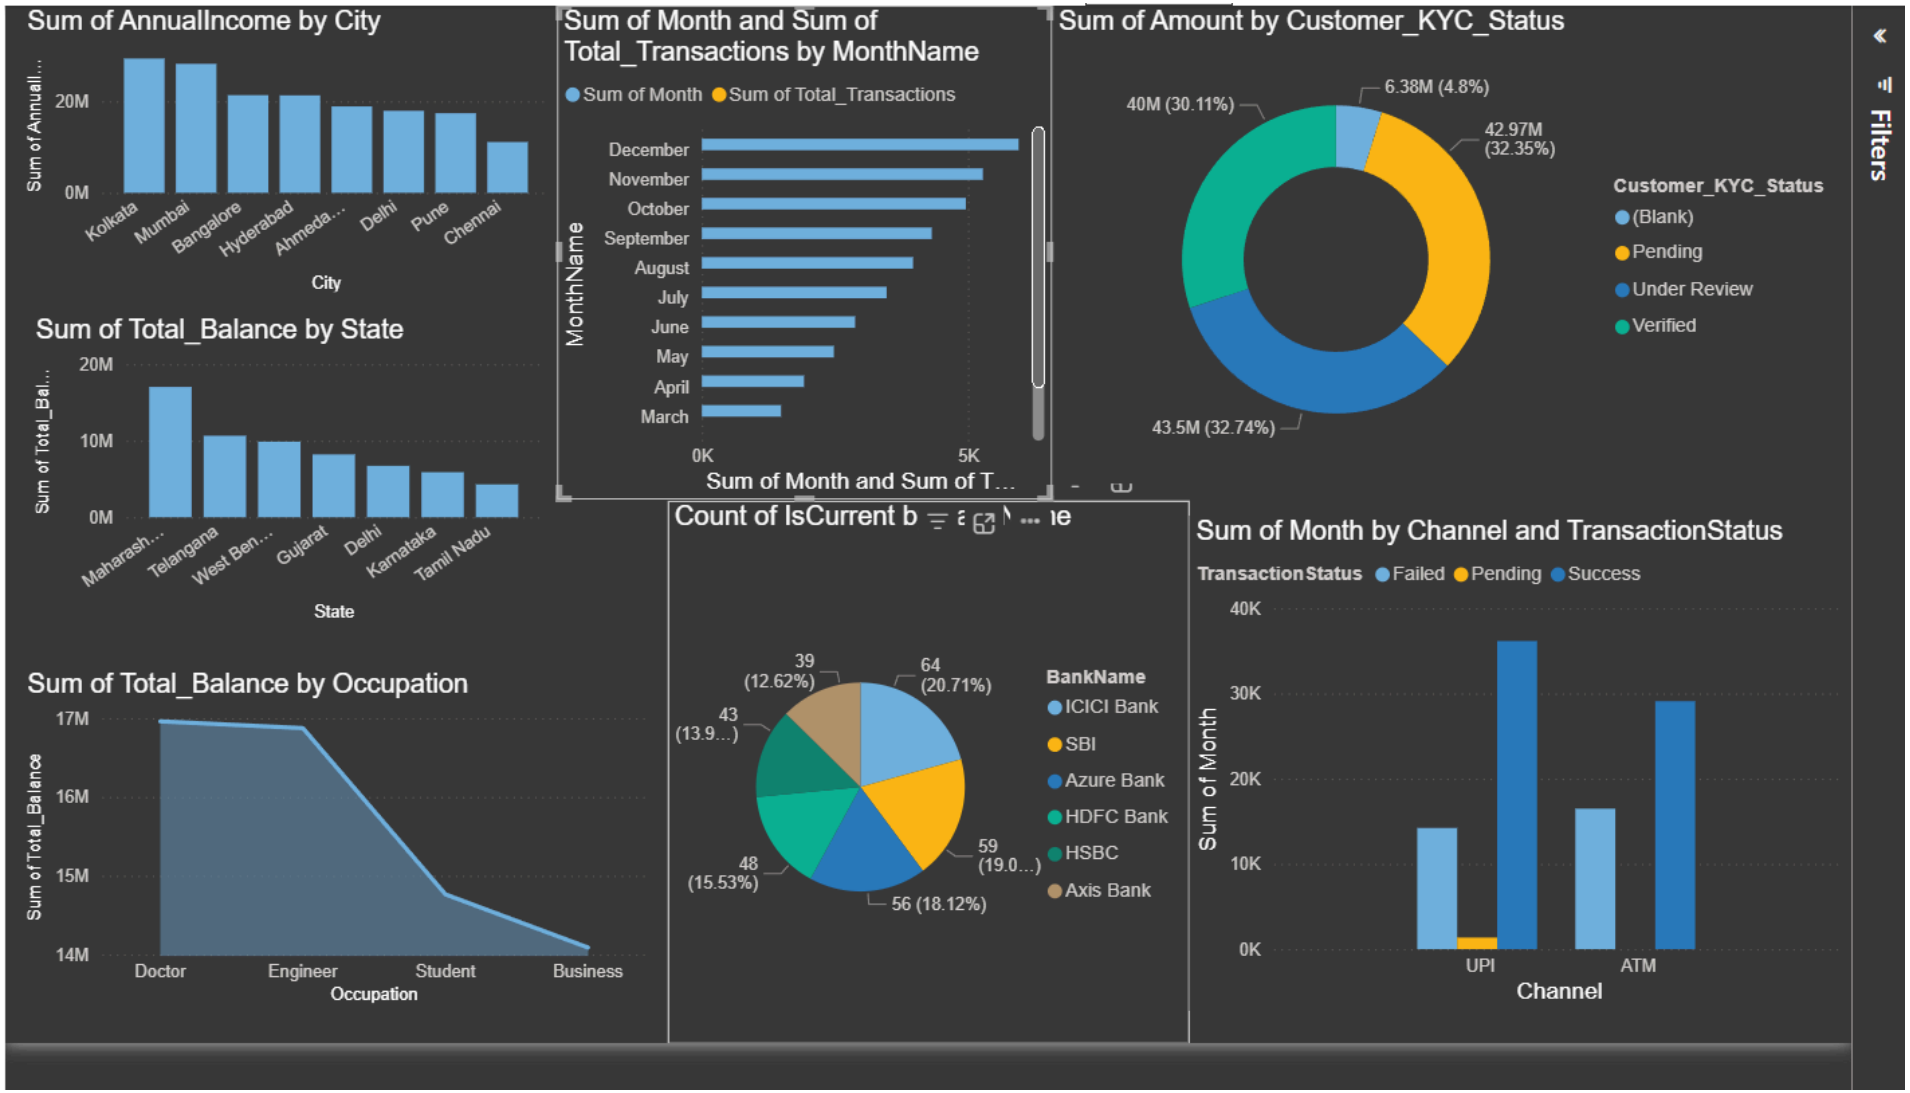
Sum of AnnualIncome by City
Sum of Annuall...
20M
0M
Kolkata
Mumbai
Bangalore
Hyderabad
Ahmeda...
Delhi
Pune
Chennai
City
Sum of Total_Balance by State
Sum of Total_Bal...
20M
10M
0M
Maharash...
Telangana
West Ben...
Gujarat
Delhi
Karnataka
Tamil Nadu
State
Sum of Total_Balance by Occupation
Sum of Total_Balance
17M
16M
15M
14M
Doctor
Engineer
Student
Business
Occupation
Sum of Month and Sum of Total_Transactions by MonthName
Sum of Month
Sum of Total_Transactions
MonthName
December
November
October
September
August
July
June
May
April
March
0K
5K
Sum of Month and Sum of T...
Count of IsCurrent b
39 (12.62%)
64 (20.71%)
43 (13.9...)
59 (19.0...)
48 (15.53%)
56 (18.12%)
BankName
ICICI Bank
SBI
Azure Bank
HDFC Bank
HSBC
Axis Bank
Sum of Amount by Customer_KYC_Status
40M (30.11%)
6.38M (4.8%)
42.97M (32.35%)
43.5M (32.74%)
Customer_KYC_Status
(Blank)
Pending
Under Review
Verified
Sum of Month by Channel and TransactionStatus
TransactionStatus
Failed
Pending
Success
40K
30K
20K
10K
0K
Sum of Month
UPI
ATM
Channel
Filters